


WASMine

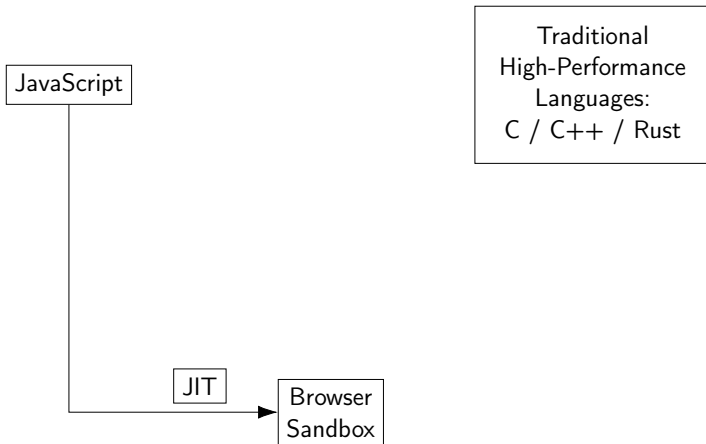
A WebAssembly Runtime with AOT, JIT and Interpreter Backends

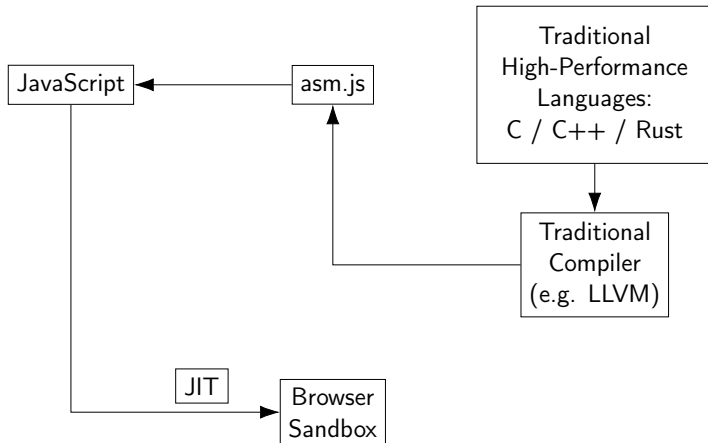
Lukas Döllner, Enrico Fazzi

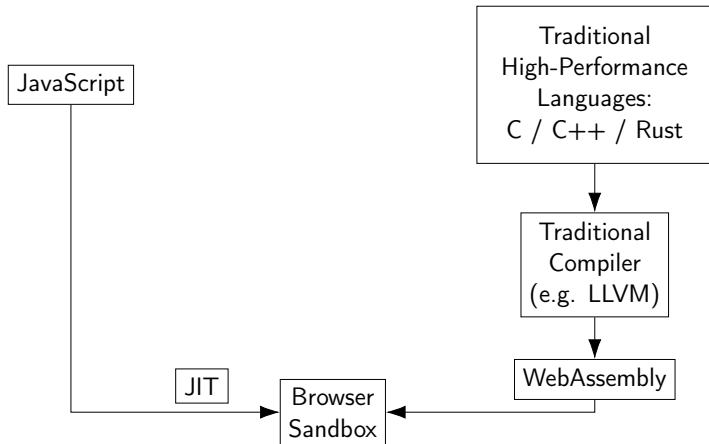
TUM School of Computation, Information, and Technology
Technical University of Munich

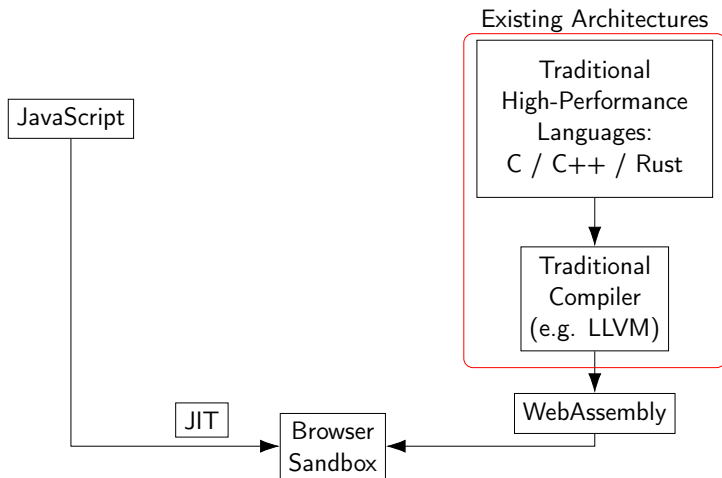
October 17, 2024

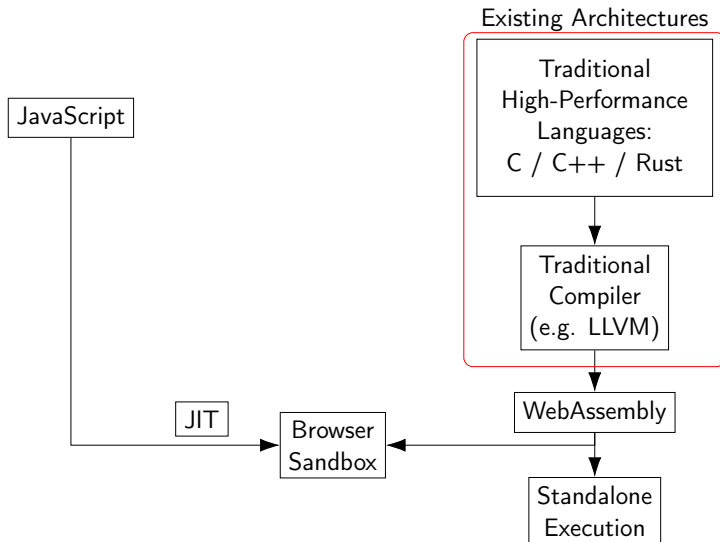












Browser JavaScript Engines:

- ▶ Google: v8
- ▶ Apple: JavaScriptCore
- ▶ Mozilla: SpiderMonkey

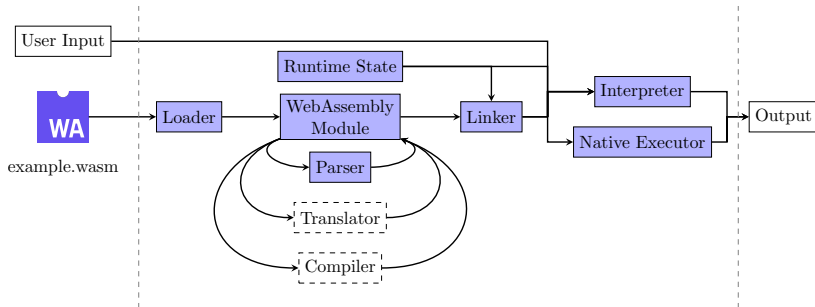
Standalone Runtimes:

- ▶ Wasmtime
- ▶ Wasmer
- ▶ WasmEdge
- ▶ ...

- ▶ **Portable:** Low system requirements
 - ▶ 64 KiB pages
 - ▶ 32-bit and 64-bit integers and IEEE 754 floats
 - ▶ Memory abstraction: Linear array of 32-bit addressable bytes


- ▶ **Portable:** Low system requirements
 - ▶ 64 KiB pages
 - ▶ 32-bit and 64-bit integers and IEEE 754 floats
 - ▶ Memory abstraction: Linear array of 32-bit addressable bytes
- ▶ **Simple:** Minimal instruction set
 - ▶ Stack-machine-based instructions
 - ▶ Simple arithmetic
 - ▶ load & store for memory accesses
 - ▶ `memory.grow n` to increase memory size
 - ▶ ...


- ▶ **Portable:** Low system requirements
 - ▶ 64 KiB pages
 - ▶ 32-bit and 64-bit integers and IEEE 754 floats
 - ▶ Memory abstraction: Linear array of 32-bit addressable bytes
- ▶ **Simple:** Minimal instruction set
 - ▶ Stack-machine-based instructions
 - ▶ Simple arithmetic
 - ▶ load & store for memory accesses
 - ▶ `memory.grow n` to increase memory size
 - ▶ ...
- ▶ **Safe:** Designed for sandboxed execution
 - ▶ No direct access to the host system
 - ▶ Memory access is bounds checked





- ▶ Numerically encoded Function Types for efficient storage and comparison
- ▶ Direct Wasm bytecode to LLVM IR translation (single pass)
- ▶ LLVM AOT compilation (+ object file loading)
- ▶ Lazy function table population
- ▶ Lazy function (code) loading

- ▶ SPEC CPU[®],2017
 - ▶ Computationally intensive, real-world programs
 - ▶ Only the 505.mcf_r, 508.namd_r, 519.lbm_r, 557.xz_r are compilable to WebAssembly (out of the box)
- ▶ PolyBench
 - ▶ 30 numerical computations, extracted from real-world applications
 - ▶ Commonly used for benchmarking compilers
 - ▶ Free and Open Source




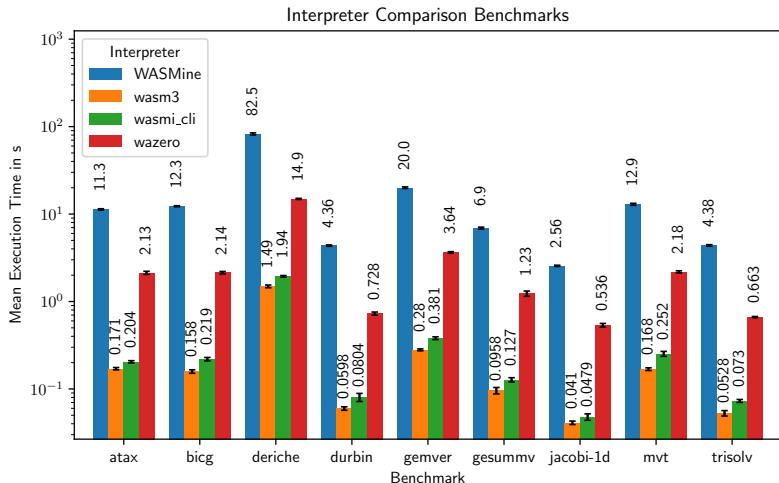




Benchmarks

PolyBench AOT Compilation Time





```
1 int main() {
2     int i, t1 = 0, t2 = 1, nextTerm, n = 10;
3
4     puts("Fibonacci Series: ");
5     for (i = 1; i <= n; ++i) {
6         printf("%d", t1);
7         nextTerm = t1 + t2;
8         t1 = t2;
9         t2 = nextTerm;
10
11         if(i < n) {
12             printf(", ");
13         }
14     }
15     puts("");
16     return 0;
17 }
```

```
1 ~/.local/share/wasi-sdk/bin/clang ./fibonacci.c -o fibonacci.wasm
```

LLVM JIT

```
1 wasm_rt -b llvm run-wasi ./fibonacci.wasm
```

LLVM AOT

```
1 wasm_rt -b llvm compile ./fibonacci.wasm -o ./fibonacci.cwasm  
2 wasm_rt -b llvm run-wasi ./fibonacci.cwasm
```

Interpreter

```
1 wasm_rt -b interpreter run-wasi ./fibonacci.wasm
```

- ▶ **Optimizations:** Further optimize parsing, de-/serialization (IR), memory management (caching), and interpreter
- ▶ **Testing:** Interpreting an interpreter
- ▶ **Documentation:** ~~Improve~~ Write the documentation

During this course we have...

- ▶ implemented a WebAssembly runtime with AOT, JIT, and Interpreter backends and benchmarked all of them
- ▶ evaluated the runtime with the WebAssembly specification tests
- ▶ benchmarked our runtime with the SPEC CPU 2017 and PolyBench benchmark suites

- ▶ WebAssembly itself
- ▶ Parser design
- ▶ LLVM's IR and (C-)API
- ▶ Interpreter design
- ▶ Cooperation is key